

50 docker commands you should know and their use cases

[Leave a Comment](#) / By [cd](#) / May 20, 2023

1. `docker run` – run a container from an image
2. `docker pull` – pull an image from a registry
3. `docker push` – push an image to a registry
4. `docker build` – build an image from a Dockerfile
5. `docker ps` – list running containers
6. `docker stop` – stop a running container
7. `docker start` – start a stopped container
8. `docker restart` – restart a container
9. `docker logs` – show the logs of a container
10. `docker exec` – execute a command inside a running container
11. `docker images` – list available images
12. `docker rm` – remove a container
13. `docker rmi` – remove an image
14. `docker inspect` – show information about a container
15. `docker network create` – create a network for containers to communicate
16. `docker network connect` – connect a container to a network
17. `docker network disconnect` – disconnect a container from a network
18. `docker port` – show the mapped ports of a container
19. `docker cp` – copy files between a container and the host
20. `docker commit` – create a new image from a container's changes
21. `docker login` – log in to a registry
22. `docker logout` – log out of a registry
23. `docker tag` – tag an image with a new name
24. `docker export` – export the contents of a container as a tar archive
25. `docker import` – create a new image from a tar archive
26. `docker save` – save an image as a tar archive
27. `docker load` – load an image from a tar archive
28. `docker top` – show the processes running inside a container

29. `docker stats` – show resource usage statistics of containers
30. `docker diff` – show the changes made to a container's filesystem
31. `docker events` – show the events generated by Docker
32. `docker history` – show the history of an image
33. `docker pause` – pause a running container
34. `docker unpause` – unpause a paused container
35. `docker kill` – send a signal to a container to stop it abruptly
36. `docker wait` – wait for a container to exit and return its exit code
37. `docker attach` – attach to a running container's console
38. `docker buildx` – build and push multi-platform images
39. `docker compose` – manage multi-container applications with Docker Compose
40. `docker swarm` – create and manage a cluster of Docker nodes
41. `docker volume create` – create a named volume for persistent data storage
42. `docker volume ls` – list available volumes
43. `docker volume rm` – remove a named volume
44. `docker system prune` – remove all unused objects from Docker
45. `docker system df` – show the usage of Docker objects
46. `docker system events` – show the events generated by Docker on the system
47. `docker system info` – show the system-wide information about Docker
48. `docker system inspect` – show detailed information about Docker objects
49. `docker system logs` – show the system logs of Docker
50. `docker system version` – show the version of Docker installed on the system

1-Which one of these frameworks are not data pipeline frameworks?

- 1-scikit
- 2-scipy
- 3-akka
- 4-apache spark
- 5-apache kafka
- 6-apache storm
- 7-don't know

The framework that is not a data pipeline framework from the list is "scikit" (option 1). Scikit-learn is a machine learning library in Python, and while it can be used as part of a larger data pipeline, it is not designed specifically for data processing or movement.

On the other hand, options 3-6, namely Akka, Apache Spark, Apache Kafka, and Apache Storm, are all frameworks that are commonly used in data pipelines for different purposes. Scipy (option 2) is a library for scientific and technical computing, including modules for optimization, signal processing, and statistics, and is also not primarily designed for data pipeline processing.

2-What are stream processing frameworks?

Note: There are 2 correct answers. Find both.

a-spark streaming

b-akka streams

c-R

d- GEvent

E- Scipy

F-don't know

The two stream processing frameworks from the list are:

a- Apache Spark Streaming: It is an extension of the Spark API that enables scalable, high-throughput, fault-tolerant processing of data streams.

b- Akka Streams: It is a set of libraries for building scalable, resilient, and reactive stream processing systems using the Akka toolkit.

R (option c) is a programming language and environment for statistical computing and graphics, and is not a stream processing framework. GEvent (option d) is a Python library for asynchronous programming and networking, and is also not a stream processing framework. Scipy (option e) is a library for scientific and technical computing, and is not primarily designed for stream processing.

3-

What is Event Sourcing?

Note: There is 1 correct answer

a-having an event source

b-saving raw events for replay in the future

c-revising events to the source, failure , propagation

d-finding the source for a specific event

E-don't know

The correct answer is:

b- Event Sourcing is a pattern in software design where the state of an application is recorded as a sequence of events rather than in its current state. In this pattern, every change to the application state is captured as an event and stored in an append-only log, also known as an event log. By replaying the events in the log, the application state can be reconstructed at any point in time, making it easier to debug, audit, and analyze the system. Event sourcing is often used in conjunction with command-query responsibility segregation (CQRS) to separate the responsibility of writing state changes from reading them.

4-What is CQRS?

Note: There is 1 correct answer

a-A system to save data in a specific structure

b- A software architecture pattern

c- an old java standard for adapters

d- a data lake approach

E- don't know

The correct answer is:

b- CQRS stands for Command Query Responsibility Segregation, which is a software architecture pattern that separates the responsibilities of handling write (command) and read (query) operations into separate components. In this pattern, a system is divided

into two parts: the command side and the query side. The command side is responsible for handling the write operations and updating the application state, while the query side is responsible for handling the read operations and serving data to the user interface or other downstream systems. CQRS helps to improve the scalability, performance, and maintainability of the system by optimizing each side for its specific responsibilities.

5-

What is a distributed log?

a-A logging system distributed for application logs

b-A distributed message pipeline with pull-pattern using offsets

c-Logging from different sources

d-A log of logs

e-An application logs across companies

f- Don't know

What is a distributed log?

a-A logging system distributed for application logs

b-A distributed message pipeline with pull-pattern using offsets

c-Logging from different sources

d-A log of logs

e-An application logs across companies

f- Don't know

The correct answer is:

b- A distributed log is a data structure that is used to manage a sequence of records that are distributed across multiple nodes in a distributed system. It is a system for storing and processing streams of records in a fault-tolerant and scalable way. A distributed log is designed to be a central source of truth for all the data that flows through a distributed system, and provides strong durability and consistency guarantees. The records in a distributed log are typically immutable and are stored

in an append-only fashion. Apache Kafka is an example of a distributed log that is commonly used in stream processing applications.

6-What of these are binary transfer protocols (choose multiple)?

Note: There is at least 1 correct answer. Find all.

- a-Thrift
- b-Hadoop
- c-Protobuf
- d-Logstash
- E-Avro
- F-don't know

The binary transfer protocols from the list are:

a- Thrift: It is a binary protocol that is used to define and serialize structured data for communication between systems. It is commonly used for RPC (remote procedure call) and is developed by Apache.

b- Protobuf: It is a language-neutral, platform-neutral binary serialization format developed by Google. It is used for efficient communication between systems and is often used for defining and sending messages in distributed systems.

c- Avro: It is a binary data serialization system developed by Apache that provides rich data structures and a compact, fast, and binary data format. It is commonly used in big data systems such as Hadoop, and provides a schema-based serialization that enables compatibility between different versions of data.

Hadoop (option b) and Logstash (option d) are not binary transfer protocols. Hadoop is a distributed data processing framework, and Logstash is an open source tool for collecting, parsing, and storing logs for searching and analyzing.

7-What is a data lake?

Note: There is at least 1 correct answer. Find all.

- a-Fast access to all the data

- b-Long-term storage for data unstructured
- c-Event sourcing pool for data
- d-The incoming data is called data lake
- e-All data out there is called data lake
- f-Don't know

The correct answers are:

b- A data lake is a centralized repository that allows storage of all structured and unstructured data at any scale. It enables the use of big data analytics, real-time analytics, and machine learning to extract insights from the stored data. Unlike a traditional data warehouse, a data lake stores data in its raw, unprocessed form and allows users to access and process it on demand. The data is not structured in advance, so it can be structured later according to the user's needs.

a- Fast access to all the data is also a characteristic of a data lake, as it provides a unified view of all the data and allows for easy access to it.

c- Event sourcing pool for data is not a correct answer. Event sourcing is a software design pattern that involves storing the state of an application as a sequence of events, rather than in its current state. A data lake is a storage system for all kinds of data, not just events.

d- The incoming data is called data lake is not a correct answer. A data lake is a storage system for data, not the data itself.

e- All data out there is called data lake is not a correct answer. A data lake is a specific storage system designed for storing large volumes of structured and unstructured data.

8-Which one of these cloud services are not used for data pipelines?

Note: There are 2 correct answers. Find both.

- a-aws EC2
- b-Firebase
- c-DigitalOcean spaces
- d-aws polly
- e-aws kinesis

f-aws dynamoDB

The two cloud services that are not used for data pipelines are:

b- Firebase: Firebase is a mobile and web application development platform developed by Google, which offers real-time databases, analytics, and authentication services. It is not typically used for building data pipelines.

c- DigitalOcean spaces: DigitalOcean Spaces is a cloud object storage service that allows users to store and retrieve data from the cloud. While it can be used for storing data that might be part of a data pipeline, it is not a service specifically designed for building data pipelines.

The correct answers are b and c.

9-What tools can you use for monitoring?

Note: There is at least 1 correct answer. Find all.

a-Icinga

b-AWS polly

c-Grafana

D-loggly

e-aws Redshift

f-don't know

The correct answers for tools that can be used for monitoring are:

a- Icinga: Icinga is an open-source monitoring tool that can monitor network resources, servers, applications, and services. It provides alerts and notifications when issues are detected and can be integrated with other monitoring tools.

c- Grafana: Grafana is an open-source platform used for data visualization and monitoring. It can be used to create dashboards and graphs that show data from various sources, such as time series databases and log files.

d- Loggly: Loggly is a cloud-based log management and analysis tool that allows users to search, analyze, and visualize log data. It provides real-time alerts and can be used to monitor applications, servers, and systems.

The correct answers are a, c, and d.

e- AWS Redshift is a cloud-based data warehousing service and not typically used for monitoring.

Therefore, the correct answer is a, c, and d.

10-Where would you potentially store logs?

a-aws s3

b-dynamodb

c-Logstash

d-aws cloudwatch

e-grafana

f-don't know

The potential places where logs can be stored are:

a- AWS S3: Amazon S3 is a cloud-based object storage service that can be used for storing and retrieving large amounts of data, including log files.

b- DynamoDB: DynamoDB is a NoSQL database service offered by AWS that can be used for storing structured data, including log data.

c- Logstash: Logstash is an open-source tool for collecting, parsing, and storing log data. It can be used to collect log data from multiple sources and store it in various destinations, including Elasticsearch, AWS S3, and other data stores.

d- AWS CloudWatch: Amazon CloudWatch is a monitoring and observability service that can be used to monitor logs, metrics, and events. It can also be used to store log data for a limited period.

Therefore, the potential places where logs can be stored are a, b, c, and d.

e- Grafana is a platform for data visualization and monitoring and not typically used for storing log data.

The correct answers are a, b, c, and d.

What is duckdb.....?

DuckDB is an open-source, in-memory analytical SQL database management system designed to efficiently process complex queries on large datasets. It is optimized for analytical workloads and aims to provide high-performance query execution while maintaining a small memory footprint.

DuckDB is written in C++ and is compatible with SQL and ANSI SQL standards. It has a modular architecture that allows users to extend its functionality using plugins. It is also designed to work seamlessly with popular programming languages such as Python, R, and Julia, making it a popular choice among data scientists and analysts. Additionally, DuckDB is fully ACID compliant, which ensures data integrity and consistency in the event of failures or errors.

what the different between mariadb,u can choose 1 more

MySQL,SQL Server,Oracle,Postgres,MariaDB,MongoDB,DynamoDB,MySQL,SQL Server,duckdb

The main difference between the traditional relational database management systems (RDBMS) such as MySQL, SQL Server, Oracle, Postgres, MariaDB and newer systems like MongoDB, DynamoDB, and DuckDB is the underlying data model they use to store and retrieve data.

The traditional RDBMS such as MySQL, SQL Server, Oracle, Postgres, and MariaDB use a structured data model where data is stored in tables with a fixed schema, and relationships between tables are defined by primary and foreign keys. These systems are optimized for transactional workloads where data consistency and integrity are paramount.

On the other hand, newer systems like MongoDB and DynamoDB use a document-oriented data model, where data is stored as JSON documents with flexible schema. These systems are optimized for highly scalable, distributed workloads, and are often used in modern web and mobile applications.

DuckDB, as mentioned earlier, is an in-memory analytical SQL database management system that is designed to efficiently process complex queries on large datasets. It is optimized for analytical workloads and aims to provide high-performance query execution while maintaining a small memory footprint. Unlike traditional RDBMS and newer systems like MongoDB and DynamoDB, DuckDB does not persist data to disk and is intended for in-memory use cases only.

In summary, the main differences between these systems are the underlying data model, intended use cases, and performance characteristics. Choosing the right system depends on the specific requirements of your application, such as data size, query complexity, concurrency, and scalability.

DuckDB is primarily an in-memory database, meaning it stores and processes data in memory for fast access and query execution. By keeping data in memory, DuckDB can provide high performance for analytical workloads, as memory access is significantly faster than accessing data from the hard disk.

However, DuckDB also supports optional persistence to disk, allowing data to be saved and loaded between different sessions or in case of system restarts. This persistence feature enables DuckDB to handle datasets that are larger than the available memory by swapping data in and out of disk storage as needed.

In summary, DuckDB primarily operates in memory for optimal performance but provides the flexibility to persist data to disk when necessary.

MySQL is a relational database management system (RDBMS) that can operate with both in-memory and disk-based storage depending on the configuration and setup.

By default, MySQL stores its data on the hard disk. When you create tables and insert data into them, the data is typically persisted on the disk in database files. MySQL utilizes a buffer pool in memory to cache frequently accessed data, which helps improve performance by reducing disk I/O. The buffer pool holds data pages that are frequently accessed, allowing faster retrieval and modification of data.

However, it's important to note that MySQL's behavior can be configured to optimize for different scenarios. For example, you can adjust the size of the buffer pool to allocate more memory for caching data in memory. This can enhance performance by minimizing disk reads and writes.

In summary, while MySQL primarily stores its data on the hard disk, it uses memory caching mechanisms to improve performance by reducing the need for frequent disk

access. The specific performance characteristics can vary based on the configuration settings and the amount of available memory.

When comparing MySQL and DuckDB, there are several factors to consider:

1. Use case: MySQL is a widely used relational database management system suitable for a broad range of applications, including transactional and analytical workloads. DuckDB, on the other hand, is specifically designed for analytical querying and excels at executing complex analytical queries efficiently.
2. Performance: Due to its focus on analytical querying and columnar storage, DuckDB is often faster than MySQL for analytical workloads. DuckDB's in-memory processing and optimized columnar storage can provide significant performance benefits for analytical queries. However, for transactional workloads or scenarios where frequent read and write operations are required, MySQL's indexing and caching mechanisms can be beneficial.
3. Data volume: MySQL is well-suited for managing large volumes of data, even beyond the available memory, by utilizing disk storage. DuckDB is primarily an in-memory database but also supports optional persistence to disk. This means DuckDB may be limited by the available memory for very large datasets, while MySQL can handle larger datasets by leveraging disk storage.
4. Ecosystem and maturity: MySQL has been around for a long time and has a large and mature ecosystem with extensive community support, various tools, and integration options. DuckDB is a relatively newer database with a smaller ecosystem but is gaining attention and growing community support.
5. SQL compatibility: Both MySQL and DuckDB support SQL as their query language. However, there may be differences in the specific SQL syntax, supported functions, and extensions. MySQL is a more traditional RDBMS, while DuckDB focuses on providing advanced analytical SQL capabilities.

Ultimately, the choice between MySQL and DuckDB depends on your specific requirements and use case. If you primarily need a versatile RDBMS for a wide range of workloads, MySQL is a solid choice. If you prioritize fast analytical querying and have relatively smaller datasets that can fit in memory, DuckDB may offer better performance for your analytical workload.

s3fs is a Python library that provides a file system interface to Amazon S3 (Simple Storage Service).

Amazon S3 is a scalable cloud storage service offered by Amazon Web Services (AWS), which allows you to store and retrieve large amounts of data.

The s3fs library allows you to mount S3 buckets as a file system, enabling you to interact with the objects in S3 using familiar file system operations such as reading, writing, and deleting files. It provides an abstraction layer that makes it easier to work with S3 storage from within Python applications.

With s3fs, you can perform operations like listing the contents of an S3 bucket, uploading and downloading files, creating directories, and more. It handles the necessary communication and authentication with the S3 service, making it convenient to integrate S3 storage into your applications.

s3fs is commonly used in various scenarios, such as data processing pipelines, backup and restore operations, distributed computing frameworks, and any application that needs to access and manipulate data stored in Amazon S3.

what the different between cluster vs server vs cloud

Clusters, servers, and the cloud are all related to computing infrastructure, but they have different characteristics and functions. Here's an overview of each term:

1. **Cluster:** A cluster refers to a group of interconnected computers or servers that work together to perform a specific task or provide a particular service. Clusters are commonly used in high-performance computing environments to achieve increased processing power, improved reliability, and fault tolerance. In a cluster, multiple individual machines are connected and treated as a single system, enabling parallel processing and workload distribution across the cluster nodes.

2. **Server:** A server is a computer or a system that provides resources, services, or functionality to other computers or users on a network. Servers are designed to handle specific tasks, such as hosting websites, managing databases, storing and sharing files, running applications, or providing network services like email or DNS. They are typically more powerful and have more processing, storage, and memory capabilities compared to regular desktop computers, allowing them to handle larger workloads and serve multiple clients simultaneously.
3. **Cloud:** The cloud refers to a network of servers, storage systems, and other computing resources that are accessed remotely over the internet. It allows users to access and utilize computing resources on-demand, without the need for physical infrastructure or hardware maintenance. Cloud computing offers scalability, flexibility, and cost-effectiveness as users can provision and pay for resources as needed. It includes services such as virtual machines (VMs), storage, databases, networking, and various software applications that are hosted and managed by cloud service providers. Cloud computing can be deployed as public, private, or hybrid clouds, depending on the specific needs and requirements of an organization.

In summary, a cluster is a group of interconnected computers working together, a server is a single computer providing services, and the cloud is a network of remote servers providing scalable and on-demand computing resources.